

Telecom_Churn_Final

October 14, 2024

0.1 Telecom Churn Case Study

With 21 predictor variables we need to predict whether a particular customer will switch to another telecom provider or not. In telecom terminology, this is referred to as churning and not churning, respectively.

```
[205]: # Suppressing Warnings
import warnings
warnings.filterwarnings('ignore')
```

0.1.1 Step 1: Importing and Merging Data

```
[206]: # Importing Pandas and NumPy
import pandas as pd, numpy as np
```

```
[207]: # Importing all datasets
churn_data = pd.read_csv("churn_data.csv")
churn_data.head()
```

```
[207]:
```

	customerID	tenure	PhoneService	Contract	PaperlessBilling	\
0	7590-VHVEG	1	No	Month-to-month	Yes	
1	5575-GNVDE	34	Yes	One year	No	
2	3668-QPYBK	2	Yes	Month-to-month	Yes	
3	7795-CFOCW	45	No	One year	No	
4	9237-HQITU	2	Yes	Month-to-month	Yes	

	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	Electronic check	29.85	29.85	No
1	Mailed check	56.95	1889.5	No
2	Mailed check	53.85	108.15	Yes
3	Bank transfer (automatic)	42.30	1840.75	No
4	Electronic check	70.70	151.65	Yes

```
[208]: customer_data = pd.read_csv("customer_data.csv")
customer_data.head()
```

```
[208]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents
0	7590-VHVEG	Female	0	Yes	No

1	5575-GNVDE	Male	0	No	No
2	3668-QPYBK	Male	0	No	No
3	7795-CFOCW	Male	0	No	No
4	9237-HQITU	Female	0	No	No

```
[209]: internet_data = pd.read_csv("internet_data.csv")
internet_data.head()
```

```
[209]: customerID      MultipleLines  InternetService  OnlineSecurity  OnlineBackup  \
0  7590-VHVEG  No phone service                DSL                No                Yes
1  5575-GNVDE                No                DSL                Yes                No
2  3668-QPYBK                No                DSL                Yes                Yes
3  7795-CFOCW  No phone service                DSL                Yes                No
4  9237-HQITU                No      Fiber optic                No                No

DeviceProtection  TechSupport  StreamingTV  StreamingMovies
0                No           No           No                No
1                Yes           No           No                No
2                No           No           No                No
3                Yes           Yes           No                No
4                No           No           No                No
```

Combining all data files into one consolidated dataframe

```
[210]: # Merging on 'customerID'
df_1 = pd.merge(churn_data, customer_data, how='inner', on='customerID')
```

```
[211]: # Final dataframe with all predictor variables
telecom = pd.merge(df_1, internet_data, how='inner', on='customerID')
```

0.1.2 Step 2: Inspecting the Dataframe

```
[212]: # Let's see the head of our master dataset
telecom.head()
```

```
[212]: customerID  tenure  PhoneService      Contract  PaperlessBilling  \
0  7590-VHVEG      1         No  Month-to-month                Yes
1  5575-GNVDE     34         Yes    One year                No
2  3668-QPYBK      2         Yes  Month-to-month                Yes
3  7795-CFOCW     45         No    One year                No
4  9237-HQITU      2         Yes  Month-to-month                Yes

PaymentMethod  MonthlyCharges  TotalCharges  Churn  gender  ...  \
0  Electronic check           29.85         29.85   No  Female  ...
1    Mailed check           56.95        1889.5   No   Male  ...
2    Mailed check           53.85         108.15  Yes   Male  ...
3  Bank transfer (automatic)      42.30        1840.75   No   Male  ...
```

4	Electronic check	70.70	151.65	Yes	Female	...
---	------------------	-------	--------	-----	--------	-----

	Partner	Dependents	MultipleLines	InternetService	OnlineSecurity	\
0	Yes	No	No phone service	DSL	No	
1	No	No	No	DSL	Yes	
2	No	No	No	DSL	Yes	
3	No	No	No phone service	DSL	Yes	
4	No	No	No	Fiber optic	No	

	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	StreamingMovies
0	Yes	No	No	No	No
1	No	Yes	No	No	No
2	Yes	No	No	No	No
3	No	Yes	Yes	No	No
4	No	No	No	No	No

[5 rows x 21 columns]

```
[213]: # Let's check the dimensions of the dataframe
telecom.shape
```

```
[213]: (7043, 21)
```

```
[214]: # let's look at the statistical aspects of the dataframe
telecom.describe()
```

```
[214]:
```

	tenure	MonthlyCharges	SeniorCitizen
count	7043.000000	7043.000000	7043.000000
mean	32.371149	64.761692	0.162147
std	24.559481	30.090047	0.368612
min	0.000000	18.250000	0.000000
25%	9.000000	35.500000	0.000000
50%	29.000000	70.350000	0.000000
75%	55.000000	89.850000	0.000000
max	72.000000	118.750000	1.000000

```
[215]: # Let's see the type of each column
telecom.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   tenure                7043 non-null   int64
2   PhoneService          7043 non-null   object
```

```

3   Contract          7043 non-null  object
4   PaperlessBilling  7043 non-null  object
5   PaymentMethod     7043 non-null  object
6   MonthlyCharges    7043 non-null  float64
7   TotalCharges      7043 non-null  object
8   Churn             7043 non-null  object
9   gender            7043 non-null  object
10  SeniorCitizen     7043 non-null  int64
11  Partner           7043 non-null  object
12  Dependents        7043 non-null  object
13  MultipleLines     7043 non-null  object
14  InternetService   7043 non-null  object
15  OnlineSecurity    7043 non-null  object
16  OnlineBackup      7043 non-null  object
17  DeviceProtection  7043 non-null  object
18  TechSupport       7043 non-null  object
19  StreamingTV       7043 non-null  object
20  StreamingMovies   7043 non-null  object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB

```

0.1.3 Step 3: Data Preparation

Converting some binary variables (Yes/No) to 0/1

```

[216]: # List of variables to map

varlist = ['PhoneService', 'PaperlessBilling', 'Churn', 'Partner',
          ↪ 'Dependents']

# Defining the map function
def binary_map(x):
    return x.map({'Yes': 1, "No": 0})

# Applying the function to the telecom list
telecom[varlist] = telecom[varlist].apply(binary_map)

```

```

[217]: telecom.head()

```

```

[217]:   customerID  tenure  PhoneService  Contract  PaperlessBilling  \
0  7590-VHVEG      1           0  Month-to-month              1
1  5575-GNVDE     34           1      One year              0
2  3668-QPYBK      2           1  Month-to-month              1
3  7795-CFOCW     45           0      One year              0
4  9237-HQITU      2           1  Month-to-month              1

      PaymentMethod  MonthlyCharges  TotalCharges  Churn  gender  ...  \
0      Electronic check           29.85          29.85     0  Female  ...

```

1	Mailed check	56.95	1889.5	0	Male	...
2	Mailed check	53.85	108.15	1	Male	...
3	Bank transfer (automatic)	42.30	1840.75	0	Male	...
4	Electronic check	70.70	151.65	1	Female	...

	Partner	Dependents	MultipleLines	InternetService	OnlineSecurity	\
0	1	0	No phone service	DSL	No	
1	0	0	No	DSL	Yes	
2	0	0	No	DSL	Yes	
3	0	0	No phone service	DSL	Yes	
4	0	0	No	Fiber optic	No	

	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	
0	Yes	No	No	No	No	
1	No	Yes	No	No	No	
2	Yes	No	No	No	No	
3	No	Yes	Yes	No	No	
4	No	No	No	No	No	

[5 rows x 21 columns]

For categorical variables with multiple levels, create dummy features (one-hot encoded)

```
[218]: # Creating a dummy variable for some of the categorical variables and dropping
        ↳ the first one.
dummy1 = pd.get_dummies(telecom[['Contract', 'PaymentMethod', 'gender',
        ↳ 'InternetService']], drop_first=True)

# Adding the results to the master dataframe
telecom = pd.concat([telecom, dummy1], axis=1)
```

```
[219]: telecom.head()
```

```
[219]:
```

	customerID	tenure	PhoneService	Contract	PaperlessBilling	\
0	7590-VHVEG	1	0	Month-to-month	1	
1	5575-GNVDE	34	1	One year	0	
2	3668-QPYBK	2	1	Month-to-month	1	
3	7795-CFOCW	45	0	One year	0	
4	9237-HQITU	2	1	Month-to-month	1	

	PaymentMethod	MonthlyCharges	TotalCharges	Churn	gender	...	\
0	Electronic check	29.85	29.85	0	Female	...	
1	Mailed check	56.95	1889.5	0	Male	...	
2	Mailed check	53.85	108.15	1	Male	...	
3	Bank transfer (automatic)	42.30	1840.75	0	Male	...	
4	Electronic check	70.70	151.65	1	Female	...	

	StreamingTV	StreamingMovies	Contract_One year	Contract_Two year	\
0	No	No	False	False	
1	No	No	True	False	
2	No	No	False	False	
3	No	No	True	False	
4	No	No	False	False	

	PaymentMethod_Credit card (automatic)	PaymentMethod_Electronic check	\
0	False	True	
1	False	False	
2	False	False	
3	False	False	
4	False	True	

	PaymentMethod_Mailed check	gender_Male	InternetService_Fiber optic	\
0	False	False	False	
1	True	True	False	
2	True	True	False	
3	False	True	False	
4	False	False	True	

	InternetService_No
0	False
1	False
2	False
3	False
4	False

[5 rows x 29 columns]

```
[220]: telecom.MultipleLines.value_counts()
```

```
[220]: MultipleLines
No          3390
Yes         2971
No phone service    682
Name: count, dtype: int64
```

1 Creating dummy variables for the remaining categorical variables and dropping the level with big names.

```
[221]: # Creating dummy variables for the variable 'MultipleLines'
ml = pd.get_dummies(telecom['MultipleLines'], prefix='MultipleLines')
# Dropping MultipleLines_No phone service column
ml1 = ml.drop(['MultipleLines_No phone service'], axis=1)
```

```

#Adding the results to the master dataframe
telecom = pd.concat([telecom,m11], axis=1)

# Creating dummy variables for the variable 'OnlineSecurity'.
os = pd.get_dummies(telecom['OnlineSecurity'], prefix='OnlineSecurity')
os1 = os.drop(['OnlineSecurity_No internet service'], axis = 1)
# Adding the results to the master dataframe
telecom = pd.concat([telecom,os1], axis=1)

# Creating dummy variables for the variable 'OnlineBackup'.
ob = pd.get_dummies(telecom['OnlineBackup'], prefix='OnlineBackup')
ob1 = ob.drop(['OnlineBackup_No internet service'], axis = 1)
# Adding the results to the master dataframe
telecom = pd.concat([telecom,ob1], axis=1)

# Creating dummy variables for the variable 'DeviceProtection'.
dp = pd.get_dummies(telecom['DeviceProtection'], prefix='DeviceProtection')
dp1 = dp.drop(['DeviceProtection_No internet service'], axis = 1)
# Adding the results to the master dataframe
telecom = pd.concat([telecom,dp1], axis=1)

# Creating dummy variables for the variable 'TechSupport'.
ts = pd.get_dummies(telecom['TechSupport'], prefix='TechSupport')
ts1 = ts.drop(['TechSupport_No internet service'], axis = 1)
# Adding the results to the master dataframe
telecom = pd.concat([telecom,ts1], axis=1)

# Creating dummy variables for the variable 'StreamingTV'.
st =pd.get_dummies(telecom['StreamingTV'], prefix='StreamingTV')
st1 = st.drop(['StreamingTV_No internet service'], axis = 1)
# Adding the results to the master dataframe
telecom = pd.concat([telecom,st1], axis=1)

# Creating dummy variables for the variable 'StreamingMovies'.
sm = pd.get_dummies(telecom['StreamingMovies'], prefix='StreamingMovies')
sm1 = sm.drop(['StreamingMovies_No internet service'], axis = 1)
# Adding the results to the master dataframe
telecom = pd.concat([telecom,sm1], axis=1)

```

```
[222]: telecom.head()
```

```

[222]:   customerID  tenure  PhoneService  Contract  PaperlessBilling  \
0  7590-VHVEG      1           0  Month-to-month              1
1  5575-GNVDE     34           1      One year              0
2  3668-QPYBK      2           1  Month-to-month              1
3  7795-CFOCW     45           0      One year              0
4  9237-HQITU      2           1  Month-to-month              1

```

	PaymentMethod	MonthlyCharges	TotalCharges	Churn	gender	...	\
0	Electronic check	29.85	29.85	0	Female	...	
1	Mailed check	56.95	1889.5	0	Male	...	
2	Mailed check	53.85	108.15	1	Male	...	
3	Bank transfer (automatic)	42.30	1840.75	0	Male	...	
4	Electronic check	70.70	151.65	1	Female	...	

	OnlineBackup_No	OnlineBackup_Yes	DeviceProtection_No	\
0	False	True	True	
1	True	False	False	
2	False	True	True	
3	True	False	False	
4	True	False	True	

	DeviceProtection_Yes	TechSupport_No	TechSupport_Yes	StreamingTV_No	\
0	False	True	False	True	
1	True	True	False	True	
2	False	True	False	True	
3	True	False	True	True	
4	False	True	False	True	

	StreamingTV_Yes	StreamingMovies_No	StreamingMovies_Yes
0	False	True	False
1	False	True	False
2	False	True	False
3	False	True	False
4	False	True	False

[5 rows x 43 columns]

Dropping the repeated variables

```
[223]: # We have created dummies for the below variables, so we can drop them
telecom = telecom.
↳ drop(['Contract', 'PaymentMethod', 'gender', 'MultipleLines', 'InternetService',
↳ 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
↳ 'TechSupport', 'StreamingTV', 'StreamingMovies'], axis = 1)
```

```
[224]: #The varaible was imported as a string we need to convert it to float
telecom['TotalCharges'] = pd.to_numeric(telecom['TotalCharges'],
↳ errors='coerce')
#telecom['TotalCharges'] = telecom['TotalCharges'].
↳ convert_objects(convert_numeric=True)
```

```
[225]: telecom.info()
```



```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 32 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   customerID                               7043 non-null   object
1   tenure                                   7043 non-null   int64
2   PhoneService                             7043 non-null   int64
3   PaperlessBilling                         7043 non-null   int64
4   MonthlyCharges                           7043 non-null   float64
5   TotalCharges                             7032 non-null   float64
6   Churn                                    7043 non-null   int64
7   SeniorCitizen                            7043 non-null   int64
8   Partner                                  7043 non-null   int64
9   Dependents                               7043 non-null   int64
10  Contract_One year                        7043 non-null   bool
11  Contract_Two year                        7043 non-null   bool
12  PaymentMethod_Credit card (automatic)    7043 non-null   bool
13  PaymentMethod_Electronic check           7043 non-null   bool
14  PaymentMethod_Mailed check               7043 non-null   bool
15  gender_Male                              7043 non-null   bool
16  InternetService_Fiber optic              7043 non-null   bool
17  InternetService_No                       7043 non-null   bool
18  MultipleLines_No                         7043 non-null   bool
19  MultipleLines_Yes                       7043 non-null   bool
20  OnlineSecurity_No                       7043 non-null   bool
21  OnlineSecurity_Yes                       7043 non-null   bool
22  OnlineBackup_No                          7043 non-null   bool
23  OnlineBackup_Yes                        7043 non-null   bool
24  DeviceProtection_No                     7043 non-null   bool
25  DeviceProtection_Yes                    7043 non-null   bool
26  TechSupport_No                          7043 non-null   bool
27  TechSupport_Yes                         7043 non-null   bool
28  StreamingTV_No                          7043 non-null   bool
29  StreamingTV_Yes                         7043 non-null   bool
30  StreamingMovies_No                      7043 non-null   bool
31  StreamingMovies_Yes                     7043 non-null   bool
dtypes: bool(22), float64(2), int64(7), object(1)
memory usage: 701.7+ KB

```

Checking for Outliers

```

[226]: # Checking for outliers in the continuous variables
num_telecom = □
↳ telecom[['tenure', 'MonthlyCharges', 'SeniorCitizen', 'TotalCharges']]

```

```

[227]: # Checking outliers at 25%, 50%, 75%, 90%, 95% and 99%
num_telecom.describe(percentiles=[.25, .5, .75, .90, .95, .99])

```

```
[227]:
```

	tenure	MonthlyCharges	SeniorCitizen	TotalCharges
count	7043.000000	7043.000000	7043.000000	7032.000000
mean	32.371149	64.761692	0.162147	2283.300441
std	24.559481	30.090047	0.368612	2266.771362
min	0.000000	18.250000	0.000000	18.800000
25%	9.000000	35.500000	0.000000	401.450000
50%	29.000000	70.350000	0.000000	1397.475000
75%	55.000000	89.850000	0.000000	3794.737500
90%	69.000000	102.600000	1.000000	5976.640000
95%	72.000000	107.400000	1.000000	6923.590000
99%	72.000000	114.729000	1.000000	8039.883000
max	72.000000	118.750000	1.000000	8684.800000

Checking for Missing Values and Inputing Them

```
[228]: # Adding up the missing values (column-wise)
telecom.isnull().sum()
```

```
[228]: customerID      0
tenure                0
PhoneService          0
PaperlessBilling      0
MonthlyCharges        0
TotalCharges          11
Churn                 0
SeniorCitizen         0
Partner               0
Dependents            0
Contract_One year    0
Contract_Two year    0
PaymentMethod_Credit card (automatic) 0
PaymentMethod_Electronic check         0
PaymentMethod_Mailed check             0
gender_Male              0
InternetService_Fiber optic            0
InternetService_No       0
MultipleLines_No         0
MultipleLines_Yes        0
OnlineSecurity_No        0
OnlineSecurity_Yes       0
OnlineBackup_No          0
OnlineBackup_Yes         0
DeviceProtection_No      0
DeviceProtection_Yes     0
TechSupport_No           0
TechSupport_Yes          0
StreamingTV_No           0
```

```
StreamingTV_Yes          0
StreamingMovies_No       0
StreamingMovies_Yes      0
dtype: int64
```

It means that $11/7043 = 0.001561834$ i.e 0.1%, best is to remove these observations from the analysis

```
[229]: # Checking the percentage of missing values
round(100*(telecom.isnull().sum()/len(telecom.index)), 2)
```

```
[229]: customerID          0.00
tenure                  0.00
PhoneService           0.00
PaperlessBilling       0.00
MonthlyCharges         0.00
TotalCharges           0.16
Churn                  0.00
SeniorCitizen          0.00
Partner                0.00
Dependents             0.00
Contract_One year     0.00
Contract_Two year     0.00
PaymentMethod_Credit card (automatic) 0.00
PaymentMethod_Electronic check         0.00
PaymentMethod_Mailed check            0.00
gender_Male                 0.00
InternetService_Fiber optic          0.00
InternetService_No            0.00
MultipleLines_No             0.00
MultipleLines_Yes           0.00
OnlineSecurity_No           0.00
OnlineSecurity_Yes          0.00
OnlineBackup_No            0.00
OnlineBackup_Yes           0.00
DeviceProtection_No        0.00
DeviceProtection_Yes       0.00
TechSupport_No             0.00
TechSupport_Yes            0.00
StreamingTV_No             0.00
StreamingTV_Yes            0.00
StreamingMovies_No         0.00
StreamingMovies_Yes        0.00
dtype: float64
```

```
[230]: # Removing NaN TotalCharges rows
telecom = telecom[~np.isnan(telecom['TotalCharges'])]
```

```
[231]: # Checking percentage of missing values after removing the missing values
round(100*(telecom.isnull().sum()/len(telecom.index)), 2)
```

```
[231]: customerID      0.0
      tenure          0.0
      PhoneService    0.0
      PaperlessBilling 0.0
      MonthlyCharges   0.0
      TotalCharges     0.0
      Churn            0.0
      SeniorCitizen    0.0
      Partner          0.0
      Dependents       0.0
      Contract_One year 0.0
      Contract_Two year 0.0
      PaymentMethod_Credit card (automatic) 0.0
      PaymentMethod_Electronic check 0.0
      PaymentMethod_Mailed check 0.0
      gender_Male      0.0
      InternetService_Fiber optic 0.0
      InternetService_No 0.0
      MultipleLines_No 0.0
      MultipleLines_Yes 0.0
      OnlineSecurity_No 0.0
      OnlineSecurity_Yes 0.0
      OnlineBackup_No 0.0
      OnlineBackup_Yes 0.0
      DeviceProtection_No 0.0
      DeviceProtection_Yes 0.0
      TechSupport_No 0.0
      TechSupport_Yes 0.0
      StreamingTV_No 0.0
      StreamingTV_Yes 0.0
      StreamingMovies_No 0.0
      StreamingMovies_Yes 0.0
      dtype: float64
```

1.0.1 Step 4: Test-Train Split

```
[232]: from sklearn.model_selection import train_test_split
```

```
[233]: # Putting feature variable to X
X = telecom.drop(['Churn', 'customerID'], axis = 1)

# Putting the response variable to y
y = telecom['Churn']
```

```
[234]: X.head()
```

```
[234]:
```

	tenure	PhoneService	PaperlessBilling	MonthlyCharges	TotalCharges	\
0	1	0	1	29.85	29.85	
1	34	1	0	56.95	1889.50	
2	2	1	1	53.85	108.15	
3	45	0	0	42.30	1840.75	
4	2	1	1	70.70	151.65	

	SeniorCitizen	Partner	Dependents	Contract_One year	Contract_Two year	\
0	0	1	0	False	False	
1	0	0	0	True	False	
2	0	0	0	False	False	
3	0	0	0	True	False	
4	0	0	0	False	False	

...	OnlineBackup_No	OnlineBackup_Yes	DeviceProtection_No	\
0	...	False	True	True
1	...	True	False	False
2	...	False	True	True
3	...	True	False	False
4	...	True	False	True

	DeviceProtection_Yes	TechSupport_No	TechSupport_Yes	StreamingTV_No	\
0	False	True	False	True	
1	True	True	False	True	
2	False	True	False	True	
3	True	False	True	True	
4	False	True	False	True	

	StreamingTV_Yes	StreamingMovies_No	StreamingMovies_Yes
0	False	True	False
1	False	True	False
2	False	True	False
3	False	True	False
4	False	True	False

```
[5 rows x 30 columns]
```

```
[235]: y.head()
```

```
[235]:
```

0	0
1	0
2	1
3	0
4	1

Name: Churn, dtype: int64

```
[236]: # Splitting the data into train and test
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = 0.7,
↳ random_state = 100)
```

1.0.2 Step 5: Feature Scaling

```
[237]: #Feature scaling
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train[['tenure', 'MonthlyCharges', 'TotalCharges']] = scaler.
↳ fit_transform(X_train[['tenure', 'MonthlyCharges', 'TotalCharges']])
X_train.head()
```

```
[237]:      tenure  PhoneService  PaperlessBilling  MonthlyCharges  TotalCharges  \
879    0.019693           1           1      -0.338074      -0.276449
5790   0.305384           0           1      -0.464443      -0.112702
6498  -1.286319           1           1       0.581425      -0.974430
880   -0.919003           1           1       1.505913      -0.550676
2784  -1.163880           1           1       1.106854      -0.835971
```

```
      SeniorCitizen  Partner  Dependents  Contract_One year  \
879                0        0           0              False
5790               0        1           1              False
6498               0        0           0              False
880                0        0           0              False
2784               0        0           1              False
```

```
      Contract_Two year  ...  OnlineBackup_No  OnlineBackup_Yes  \
879                  False  ...              False              True
5790                 False  ...              False              True
6498                 False  ...              False              True
880                  False  ...              False              True
2784                 False  ...               True              False
```

```
      DeviceProtection_No  DeviceProtection_Yes  TechSupport_No  \
879                   True                   False              True
5790                  True                   False              True
6498                  False                  True              True
880                   False                  True              False
2784                  False                  True              False
```

```
      TechSupport_Yes  StreamingTV_No  StreamingTV_Yes  StreamingMovies_No  \
879                False              True          False              True
5790                False             False           True              False
6498                False              True          False              True
880                 True             False           True              False
2784                True             False           True              False
```

	StreamingMovies_Yes
879	False
5790	True
6498	False
880	True
2784	True

[5 rows x 30 columns]

```
[238]: ### Checking the Churn Rate
churn = (sum(telecom['Churn']) / len(telecom['Churn'].index)) * 100
churn


# We have almost 26.5% of churn rate
```

[238]: 26.578498293515356

1.0.3 Step 6: Looking at Correlations

```
[239]: import matplotlib.pyplot as plt
        %matplotlib inline
        import seaborn as sns
```

```
[240]: # let us see the correlation matrix
plt.figure(figsize=(20,10))
sns.heatmap(X_train.corr(), annot = True)
plt.show()
```




[241]: *#### Dropping highly correlated dummy variables*

```
X_train = X_train.drop(['MultipleLines_No', 'OnlineSecurity_No',
↳ 'OnlineBackup_No',
↳ 'DeviceProtection_No', 'TechSupport_No', 'StreamingTV_No', 'StreamingMovies_No'],
↳ axis = 1)
X_test = X_test.drop(['MultipleLines_No', 'OnlineSecurity_No',
↳ 'OnlineBackup_No',
↳ 'DeviceProtection_No', 'TechSupport_No', 'StreamingTV_No', 'StreamingMovies_No'],
↳ axis = 1)
```

[242]: *# let us see the correlation matrix after dropping highly correlated columns*

```
plt.figure(figsize=(20,10))
sns.heatmap(X_train.corr(), annot = True)
plt.show()
```

```
[243]: X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 4922 entries, 879 to 5649
```

```
Data columns (total 23 columns):
```

#	Column	Non-Null Count	Dtype
0	tenure	4922 non-null	float64
1	PhoneService	4922 non-null	int64
2	PaperlessBilling	4922 non-null	int64
3	MonthlyCharges	4922 non-null	float64
4	TotalCharges	4922 non-null	float64
5	SeniorCitizen	4922 non-null	int64
6	Partner	4922 non-null	int64
7	Dependents	4922 non-null	int64
8	Contract_One year	4922 non-null	bool
9	Contract_Two year	4922 non-null	bool
10	PaymentMethod_Credit card (automatic)	4922 non-null	bool
11	PaymentMethod_Electronic check	4922 non-null	bool
12	PaymentMethod_Mailed check	4922 non-null	bool
13	gender_Male	4922 non-null	bool
14	InternetService_Fiber optic	4922 non-null	bool
15	InternetService_No	4922 non-null	bool
16	MultipleLines_Yes	4922 non-null	bool
17	OnlineSecurity_Yes	4922 non-null	bool

```

18 OnlineBackup_Yes          4922 non-null    bool
19 DeviceProtection_Yes     4922 non-null    bool
20 TechSupport_Yes          4922 non-null    bool
21 StreamingTV_Yes          4922 non-null    bool
22 StreamingMovies_Yes      4922 non-null    bool
dtypes: bool(15), float64(3), int64(5)
memory usage: 418.2 KB

```

```

[244]: bool_columns = X_train.select_dtypes(include=['bool']).columns
X_train[bool_columns] = X_train[bool_columns].astype(int)

```

1.0.4 Step 7: Model Building

Let's start by splitting our data into a training set and a test set.

```

[245]: # Running Your First Training Model
import statsmodels.api as sm

# Logistic regression model
X_train_const = sm.add_constant(X_train)
logm1 = sm.GLM(y_train, X_train_const, family=sm.families.Binomial())
result = logm1.fit()
print(result.summary())

```

Generalized Linear Model Regression Results

```

=====
Dep. Variable:          Churn    No. Observations:          4922
Model:                GLM      Df Residuals:              4898
Model Family:         Binomial  Df Model:                  23
Link Function:         Logit    Scale:                    1.0000
Method:               IRLS     Log-Likelihood:          -2004.7
Date:                 Wed, 07 Aug 2024    Deviance:                 4009.4
Time:                 03:42:23    Pearson chi2:             6.07e+03
No. Iterations:         7        Pseudo R-squ. (CS):       0.2844
Covariance Type:       nonrobust
=====
=====

```

			coef	std err	z
P> z	[0.025	0.975]			
const			-3.9382	1.546	-2.547
0.011	-6.969	-0.908			
tenure			-1.5172	0.189	-8.015
0.000	-1.888	-1.146			
PhoneService			0.9507	0.789	1.205
0.228	-0.595	2.497			
PaperlessBilling			0.3254	0.090	3.614

0.000	0.149	0.502			
MonthlyCharges			-2.1806	1.160	-1.880
0.060	-4.454	0.092			
TotalCharges			0.7332	0.198	3.705
0.000	0.345	1.121			
SeniorCitizen			0.3984	0.102	3.924
0.000	0.199	0.597			
Partner			0.0374	0.094	0.399
0.690	-0.146	0.221			
Dependents			-0.1430	0.107	-1.332
0.183	-0.353	0.067			
Contract_One year			-0.6578	0.129	-5.106
0.000	-0.910	-0.405			
Contract_Two year			-1.2455	0.212	-5.874
0.000	-1.661	-0.830			
PaymentMethod_Credit card (automatic)			-0.2577	0.137	-1.883
0.060	-0.526	0.011			
PaymentMethod_Electronic check			0.1615	0.113	1.434
0.152	-0.059	0.382			
PaymentMethod_Mailed check			-0.2536	0.137	-1.845
0.065	-0.523	0.016			
gender_Male			-0.0346	0.078	-0.442
0.658	-0.188	0.119			
InternetService_Fiber optic			2.5124	0.967	2.599
0.009	0.618	4.407			
InternetService_No			-2.7792	0.982	-2.831
0.005	-4.703	-0.855			
MultipleLines_Yes			0.5623	0.214	2.628
0.009	0.143	0.982			
OnlineSecurity_Yes			-0.0245	0.216	-0.113
0.910	-0.448	0.399			
OnlineBackup_Yes			0.1740	0.212	0.822
0.411	-0.241	0.589			
DeviceProtection_Yes			0.3229	0.215	1.501
0.133	-0.099	0.744			
TechSupport_Yes			-0.0305	0.216	-0.141
0.888	-0.455	0.394			
StreamingTV_Yes			0.9598	0.396	2.423
0.015	0.183	1.736			
StreamingMovies_Yes			0.8484	0.396	2.143
0.032	0.072	1.624			

=====

=====

1.0.5 Step 8: Feature Selection Using RFE

```
[246]: # Feature selection
from sklearn.linear_model import LogisticRegression
from sklearn.feature_selection import RFE

logreg = LogisticRegression(max_iter=1000)

rfe = RFE(estimator=logreg, n_features_to_select=13)
rfe = rfe.fit(X_train, y_train)
```

```
[247]: print("Selected features:", X_train.columns[rfe.support_])
```

```
Selected features: Index(['tenure', 'MonthlyCharges', 'TotalCharges',
                          'SeniorCitizen',
                          'Contract_One year', 'Contract_Two year',
                          'PaymentMethod_Credit card (automatic)', 'PaymentMethod_Mailed check',
                          'InternetService_Fiber optic', 'InternetService_No',
                          'MultipleLines_Yes', 'StreamingTV_Yes', 'StreamingMovies_Yes'],
                          dtype='object')
```

```
[248]: rfe.support_
```

```
[248]: array([ True, False, False,  True,  True,  True, False, False,  True,
         True,  True, False,  True, False,  True,  True,  True, False,
        False, False, False,  True,  True])
```

```
[249]: list(zip(X_train.columns, rfe.support_, rfe.ranking_))
```

```
[249]: [('tenure', True, 1),
        ('PhoneService', False, 5),
        ('PaperlessBilling', False, 2),
        ('MonthlyCharges', True, 1),
        ('TotalCharges', True, 1),
        ('SeniorCitizen', True, 1),
        ('Partner', False, 9),
        ('Dependents', False, 8),
        ('Contract_One year', True, 1),
        ('Contract_Two year', True, 1),
        ('PaymentMethod_Credit card (automatic)', True, 1),
        ('PaymentMethod_Electronic check', False, 6),
        ('PaymentMethod_Mailed check', True, 1),
        ('gender_Male', False, 10),
        ('InternetService_Fiber optic', True, 1),
        ('InternetService_No', True, 1),
        ('MultipleLines_Yes', True, 1),
        ('OnlineSecurity_Yes', False, 4),
        ('OnlineBackup_Yes', False, 7),
```

```
('DeviceProtection_Yes', False, 11),
('TechSupport_Yes', False, 3),
('StreamingTV_Yes', True, 1),
('StreamingMovies_Yes', True, 1)]
```

```
[250]: col = X_train.columns[rfe.support_]
X_train.columns[~rfe.support_]
```

```
[250]: Index(['PhoneService', 'PaperlessBilling', 'Partner', 'Dependents',
'PaymentMethod_Electronic check', 'gender_Male', 'OnlineSecurity_Yes',
'OnlineBackup_Yes', 'DeviceProtection_Yes', 'TechSupport_Yes'],
dtype='object')
```

Assessing the model with StatsModels

```
[251]: X_train_sm = sm.add_constant(X_train[col])
logm2 = sm.GLM(y_train, X_train_sm, family = sm.families.Binomial())
res = logm2.fit()
print(res.summary())
```

```

                                Generalized Linear Model Regression Results
=====
Dep. Variable:                  Churn    No. Observations:                  4922
Model:                          GLM      Df Residuals:                  4908
Model Family:                   Binomial  Df Model:                      13
Link Function:                  Logit     Scale:                        1.0000
Method:                         IRLS     Log-Likelihood:               -2020.8
Date:                           Wed, 07 Aug 2024    Deviance:                     4041.6
Time:                           03:42:24    Pearson chi2:                 6.15e+03
No. Iterations:                  7        Pseudo R-squ. (CS):           0.2797
Covariance Type:                nonrobust
=====
=====
                                coef    std err          z
P>|z|      [0.025    0.975]
-----
const                                -2.2218    0.164   -13.559
0.000     -2.543    -1.901
tenure                               -1.5388    0.186    -8.257
0.000     -1.904    -1.174
MonthlyCharges                       -1.1395    0.185    -6.173
0.000     -1.501    -0.778
TotalCharges                          0.7223    0.197     3.673
0.000      0.337     1.108
SeniorCitizen                        0.4614    0.099     4.655
0.000      0.267     0.656
Contract_One year                    -0.7326    0.127    -5.769

```

0.000	-0.981	-0.484			
Contract_Two year			-1.4007	0.208	-6.722
0.000	-1.809	-0.992			
PaymentMethod_Credit card (automatic)			-0.3790	0.112	-3.376
0.001	-0.599	-0.159			
PaymentMethod_Mailed check			-0.4083	0.111	-3.690
0.000	-0.625	-0.191			
InternetService_Fiber optic			1.8340	0.198	9.276
0.000	1.446	2.221			
InternetService_No			-1.8156	0.213	-8.533
0.000	-2.233	-1.399			
MultipleLines_Yes			0.4351	0.102	4.268
0.000	0.235	0.635			
StreamingTV_Yes			0.6440	0.111	5.776
0.000	0.425	0.863			
StreamingMovies_Yes			0.5260	0.109	4.806
0.000	0.311	0.740			

=====

=====

```
[252]: # Getting the predicted values on the train set
y_train_pred = res.predict(X_train_sm)
y_train_pred[:10]
```

```
[252]: 879      0.163642
5790     0.254667
6498     0.556098
880      0.520664
2784     0.670002
3874     0.366496
5387     0.544218
6623     0.792773
4465     0.201923
5364     0.476004
dtype: float64
```

```
[253]: y_train_pred = y_train_pred.values.reshape(-1)
y_train_pred[:10]
```

```
[253]: array([0.16364211, 0.25466655, 0.55609771, 0.52066367, 0.67000227,
        0.366496   , 0.54421765, 0.7927729  , 0.2019225  , 0.47600412])
```

```
[254]: ##### Creating a dataframe with the actual churn flag and the predicted
        probabilities
y_train_pred_final = pd.DataFrame({'Churn': y_train.values, 'Churn_Prob':
        y_train_pred})
y_train_pred_final['CustID'] = y_train.index
```

```
y_train_pred_final.head()
```

```
[254]:
```

	Churn	Churn_Prob	CustID
0	0	0.163642	879
1	0	0.254667	5790
2	1	0.556098	6498
3	1	0.520664	880
4	1	0.670002	2784

```
[255]: ##### Creating new column 'predicted' with 1 if Churn_Prob > 0.5 else 0

y_train_pred_final['predicted'] = y_train_pred_final.Churn_Prob.map(lambda x: 1
    ↳ if x > 0.5 else 0)
y_train_pred_final.head()
```

```
[255]:
```

	Churn	Churn_Prob	CustID	predicted
0	0	0.163642	879	0
1	0	0.254667	5790	0
2	1	0.556098	6498	1
3	1	0.520664	880	1
4	1	0.670002	2784	1

```
[256]: # Confusion matrix

from sklearn import metrics
confusion = metrics.confusion_matrix(y_train_pred_final.Churn,
    ↳ y_train_pred_final.predicted)
print(confusion)
```

```
[[3276  359]
 [ 596  691]]
```

```
[257]: # Let's check the overall accuracy.

print(metrics.accuracy_score(y_train_pred_final.Churn, y_train_pred_final.
    ↳ predicted))
```

```
0.8059731816334823
```

```
[258]: ##### Checking VIFs

from statsmodels.stats.outliers_influence import variance_inflation_factor

vif = pd.DataFrame()
vif['Features'] = X_train_sm[col].columns
vif['VIF'] = [variance_inflation_factor(X_train_sm[col].values, i) for i in
    ↳ range (X_train_sm[col].shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
```

```
vif = vif.sort_values(by = 'VIF', ascending = False)
vif
```

```
[258]:
```

	Features	VIF
1	MonthlyCharges	13.74
2	TotalCharges	10.36
0	tenure	7.31
9	InternetService_No	4.90
8	InternetService_Fiber optic	4.57
5	Contract_Two year	2.82
11	StreamingTV_Yes	2.64
12	StreamingMovies_Yes	2.64
10	MultipleLines_Yes	2.28
4	Contract_One year	1.73
7	PaymentMethod_Mailed check	1.67
6	PaymentMethod_Credit card (automatic)	1.43
3	SeniorCitizen	1.32

```
[259]: # Dropping column which has high VIFs
```

```
col = col.drop(['MonthlyCharges'], 1)
col
```

```
[259]: Index(['tenure', 'TotalCharges', 'SeniorCitizen', 'Contract_One year',
          'Contract_Two year', 'PaymentMethod_Credit card (automatic)',
          'PaymentMethod_Mailed check', 'InternetService_Fiber optic',
          'InternetService_No', 'MultipleLines_Yes', 'StreamingTV_Yes',
          'StreamingMovies_Yes'],
          dtype='object')
```

```
[260]: # Re-run the model using the selected variables
```

```
X_train_sm = sm.add_constant(X_train[col])
logm3 = sm.GLM(y_train, X_train_sm, family = sm.families.Binomial())
res = logm3.fit()
print(res.summary())
```

Generalized Linear Model Regression Results

```
=====
Dep. Variable:          Churn    No. Observations:          4922
Model:                  GLM      Df Residuals:              4909
Model Family:           Binomial  Df Model:                  12
Link Function:          Logit     Scale:                  1.0000
Method:                 IRLS      Log-Likelihood:         -2040.2
Date:                   Wed, 07 Aug 2024    Deviance:               4080.4
Time:                   03:42:24    Pearson chi2:           5.62e+03
No. Iterations:         7          Pseudo R-squ. (CS):      0.2740
Covariance Type:        nonrobust
```



```
=====
=====
                                coef    std err          z
P>|z|      [0.025    0.975]
-----
const                -1.5257    0.113   -13.468
0.000      -1.748      -1.304
tenure               -1.2295    0.177    -6.931
0.000      -1.577      -0.882
TotalCharges          0.3023    0.183     1.651
0.099      -0.057     0.661
SeniorCitizen         0.5105    0.099     5.177
0.000       0.317     0.704
Contract_One year    -0.8131    0.126    -6.464
0.000      -1.060      -0.567
Contract_Two year   -1.5074    0.207    -7.276
0.000      -1.913      -1.101
PaymentMethod_Credit card (automatic) -0.4026    0.112    -3.597
0.000      -0.622      -0.183
PaymentMethod_Mailed check -0.4061    0.109    -3.711
0.000      -0.621      -0.192
InternetService_Fiber optic  0.8270    0.106     7.785
0.000       0.619     1.035
InternetService_No   -0.8864    0.152    -5.823
0.000      -1.185      -0.588
MultipleLines_Yes     0.1970    0.093     2.118
0.034       0.015     0.379
StreamingTV_Yes       0.2956    0.095     3.095
0.002       0.108     0.483
StreamingMovies_Yes   0.1984    0.095     2.087
0.037       0.012     0.385
=====
=====
```

```
[261]: y_train_pred = res.predict(X_train_sm).values.reshape(-1)
y_train_pred[:10]
y_train_pred_final['Churn_Prob'] = y_train_pred
```

```
[262]: # Creating new column 'predicted' with 1 if Churn_Prob > 0.5 else 0

y_train_pred_final['predicted'] = y_train_pred_final.Churn_Prob.map(lambda x: 1
↳ if x > 0.5 else 0)
```

```
[263]: # Let's check the overall accuracy.
print(metrics.accuracy_score(y_train_pred_final.Churn, y_train_pred_final.
↳ predicted))
```

0.8021129622104836

The accuracy is still practically the same.

Let's now check the VIFs again

```
[264]: vif = pd.DataFrame()
vif['Features'] = X_train_sm[col].columns
vif['VIF'] = [variance_inflation_factor(X_train_sm[col].values, i) for i in
             range(X_train_sm[col].shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = 'VIF', ascending = False)
vif
```

```
[264]:
```

	Features	VIF
1	TotalCharges	7.12
0	tenure	6.79
4	Contract_Two year	2.75
7	InternetService_Fiber optic	2.60
11	StreamingMovies_Yes	2.52
10	StreamingTV_Yes	2.51
8	InternetService_No	2.30
9	MultipleLines_Yes	2.22
3	Contract_One year	1.66
6	PaymentMethod_Mailed check	1.57
5	PaymentMethod_Credit card (automatic)	1.39
2	SeniorCitizen	1.29

```
[265]: # Re-run the model without TotalCharges
```

```
col = col.drop(['TotalCharges'], 1)
col
```

```
[265]: Index(['tenure', 'SeniorCitizen', 'Contract_One year', 'Contract_Two year',
        'PaymentMethod_Credit card (automatic)', 'PaymentMethod_Mailed check',
        'InternetService_Fiber optic', 'InternetService_No',
        'MultipleLines_Yes', 'StreamingTV_Yes', 'StreamingMovies_Yes'],
        dtype='object')
```

```
[266]: X_train_sm = sm.add_constant(X_train[col])
logm4 = sm.GLM(y_train, X_train_sm, family = sm.families.Binomial())
res = logm4.fit()
print(res.summary())
```

Generalized Linear Model Regression Results

```
=====
Dep. Variable:          Churn    No. Observations:          4922
Model:                GLM      Df Residuals:              4910
Model Family:         Binomial  Df Model:                  11
```

```

Link Function:          Logit      Scale:          1.0000
Method:                IRLS      Log-Likelihood: -2041.6
Date:                  Wed, 07 Aug 2024      Deviance:      4083.2
Time:                  03:42:25      Pearson chi2:   5.23e+03
No. Iterations:        7      Pseudo R-squ. (CS): 0.2736
Covariance Type:      nonrobust

```

			coef	std err	z
P> z	[0.025	0.975]			

const			-1.5642	0.110	-14.183
0.000	-1.780	-1.348			
tenure			-0.9594	0.065	-14.842
0.000	-1.086	-0.833			
SeniorCitizen			0.5110	0.099	5.170
0.000	0.317	0.705			
Contract_One year			-0.8054	0.125	-6.429
0.000	-1.051	-0.560			
Contract_Two year			-1.4776	0.205	-7.194
0.000	-1.880	-1.075			
PaymentMethod_Credit card (automatic)			-0.4018	0.112	-3.586
0.000	-0.621	-0.182			
PaymentMethod_Mailed check			-0.3831	0.108	-3.531
0.000	-0.596	-0.170			
InternetService_Fiber optic			0.9038	0.095	9.467
0.000	0.717	1.091			
InternetService_No			-0.9031	0.151	-5.968
0.000	-1.200	-0.607			
MultipleLines_Yes			0.2237	0.092	2.444
0.015	0.044	0.403			
StreamingTV_Yes			0.3321	0.093	3.575
0.000	0.150	0.514			
StreamingMovies_Yes			0.2334	0.093	2.519
0.012	0.052	0.415			

```

[267]: y_train_pred = res.predict(X_train_sm).values.reshape(-1)
       y_train_pred[:10]
       y_train_pred_final['Churn_Prob'] = y_train_pred

```

```

[268]: y_train_pred_final['predicted'] = y_train_pred_final.Churn_Prob.map(lambda x: 1
       ↪ if x > 0.5 else 0)

```

```
[269]: print(metrics.accuracy_score(y_train_pred_final.Churn, y_train_pred_final.
↳predicted))
```

0.802519301097115

```
[270]: vif = pd.DataFrame()
vif['Features'] = X_train_sm[col].columns
vif['VIF'] = [variance_inflation_factor(X_train_sm[col].values, i) for i in
↳range (X_train_sm[col].shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = 'VIF', ascending = False)
vif
```

```
[270]:
```

	Features	VIF
3	Contract_Two year	2.68
10	StreamingMovies_Yes	2.46
9	StreamingTV_Yes	2.44
6	InternetService_Fiber optic	2.42
8	MultipleLines_Yes	2.19
0	tenure	1.97
7	InternetService_No	1.82
2	Contract_One year	1.65
5	PaymentMethod_Mailed check	1.57
4	PaymentMethod_Credit card (automatic)	1.39
1	SeniorCitizen	1.28

All variables have a good value of VIF. So we need not drop any more variables and we can proceed with making predictions using this model only

```
[271]: # Let's take a look at the confusion matrix again
confusion = metrics.confusion_matrix(y_train_pred_final.Churn,
↳y_train_pred_final.predicted)
print(confusion)
```

```
[[3277  358]
 [ 614  673]]
```

```
[272]: # cross checking the accuracy will help us understand the model if its changed
↳or not
print(metrics.accuracy_score(y_train_pred_final.Churn, y_train_pred_final.
↳predicted))
```

0.802519301097115

2 Actual/Predicted not_churn churn

# not_churn	3277	358
# churn	614	673

2.1 Metrics beyond simply accuracy

```
[273]: TP = confusion[1,1] # true positive
TN = confusion[0,0] # true negatives
FP = confusion[0,1] # false positives
FN = confusion[1,0] # false negatives
```

```
[274]: # Let's see the sensitivity of our logistic regression model
print("sensitivity =", TP / float(TP+FN))

# Let us calculate specificity
print("specificity = ", TN / float(TN+FP))

# Calculate false positive rate - predicting churn when customer does not have
↳churned
print("False positive rate =", FP/float(TN+FP))

# positive predictive value
print("Positive predicted rate =", TP/float(TP+FP))

# Negative predictive value
print ("Negative predicted Rate =", TN / float(TN+ FN))
```

```
sensitivity = 0.5229215229215229
specificity = 0.9015130674002751
False positive rate = 0.0984869325997249
Positive predicted rate = 0.6527643064985451
Negative predicted Rate = 0.8421999485993318
```

2.1.1 Step 9: Plotting the ROC Curve


An ROC curve demonstrates several things:

- It shows the tradeoff between sensitivity and specificity (any increase in sensitivity will be accompanied by a decrease in specificity).
- The closer the curve follows the left-hand border and then the top border of the ROC space, the more accurate the test.
- The closer the curve comes to the 45-degree diagonal of the ROC space, the less accurate the test.

```
[275]: # Plot the ROC Curve
def draw_roc(actual, probs):
    fpr, tpr, thresholds = metrics.roc_curve(actual, probs,
↳drop_intermediate=False)
    auc_score = metrics.roc_auc_score(actual, probs)
    plt.figure(figsize = (5,5))
    plt.plot(fpr, tpr, label = "ROC Curve(area = %0.2f)" % auc_score)
    plt.plot([0,1], [0,1], 'k--')
    plt.xlim([0.0, 1.0])
```

```
plt.ylim([0.0, 1.05])
plt.xlabel("False positive rate or [1- True negative rate]")
plt.ylabel("True positive rate")
plt.title("Receiver Operating Characteristic Curve")
plt.legend(loc = "lower right")
plt.show()
return
```

```
[276]: fpr, tpr, thresholds = metrics.roc_curve(y_train_pred_final.Churn,
↳ y_train_pred_final.Churn_Prob, drop_intermediate = False)
draw_roc(y_train_pred_final.Churn, y_train_pred_final.Churn_Prob)
```



2.1.2 Step 10: Finding Optimal Cutoff Point

Optimal cutoff probability is that prob where we get balanced sensitivity and specificity

```
[277]: # finding the optimal cut off point

numbers = [float(x)/10 for x in range(10)]
```

```

for i in numbers:
    y_train_pred_final[i] = y_train_pred_final.Churn_Prob.map(lambda x: 1 if x >=
↪0.5 else 0)
y_train_pred_final.head()

```

```

[277]:
   Churn  Churn_Prob  CustID  predicted  0.0  0.1  0.2  0.3  0.4  0.5  0.6  \
0      0    0.204339     879          0    0    0    0    0    0    0    0
1      0    0.215562     5790         0    0    0    0    0    0    0    0
2      1    0.639615     6498          1    1    1    1    1    1    1    1
3      1    0.687136      880          1    1    1    1    1    1    1    1
4      1    0.735303     2784          1    1    1    1    1    1    1    1

      0.7  0.8  0.9
0      0    0    0
1      0    0    0
2      1    1    1
3      1    1    1
4      1    1    1

```

```

[299]: # Now let's calculate accuracy sensitivity and specificity for various
↪probability cutoffs.

cutoff_df = pd.DataFrame(columns = ['Prob', 'Accuracy', 'Sensitivity',
↪'Specificity'])
from sklearn.metrics import confusion_matrix

# TP = confusion[1,1] # true positive
# TN = confusion[0,0] # true negatives
# FP = confusion[0,1] # false positives
# FN = confusion[1,0] # false negatives

# List of probability cutoffs to evaluate
num = [0.0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9]

for i in num:
    # Create binary predictions based on the cutoff
    y_train_pred_final['Predicted_Label'] = y_train_pred_final['Churn_Prob'].
↪apply(lambda x: 1 if x >= i else 0)

# Compute the confusion matrix
    cm1 = confusion_matrix(y_train_pred_final['Churn'],
↪y_train_pred_final['Predicted_Label'])

# Calculate accuracy, sensitivity, and specificity
    total1 = sum(sum(cm1))
    Accuracy = (cm1[0,0] + cm1[1,1]) / total1

```

```

Specificity = cm1[0,0] / (cm1[0,0] + cm1[0,1])
Sensitivity = cm1[1,1] / (cm1[1,0] + cm1[1,1])

# Store the results in the DataFrame
cutoff_df.loc[i] = [i, Accuracy, Sensitivity, Specificity]

print(cutoff_df)

```

	Prob	Accuracy	Sensitivity	Specificity
0.0	0.0	0.261479	1.000000	0.000000
0.1	0.1	0.608086	0.943279	0.489409
0.2	0.2	0.716782	0.851593	0.669051
0.3	0.3	0.763308	0.781663	0.756809
0.4	0.4	0.789313	0.644911	0.840440
0.5	0.5	0.802519	0.522922	0.901513
0.6	0.6	0.799878	0.387723	0.945805
0.7	0.7	0.773466	0.184926	0.981843
0.8	0.8	0.745429	0.031080	0.998349
0.9	0.9	0.738521	0.000000	1.000000

```

[300]: # Let's plot accuracy sensitivity and specificity for various probabilities.


cutoff_df.plot.line(x = 'Prob', y = ['Accuracy', 'Sensitivity', 'Specificity'])
plt.show

```

```

[300]: <function matplotlib.pyplot.show(close=None, block=None)>

```

From the curve above, 0.3 is the optimum point to take it as a cutoff probability.

```
[302]: y_train_pred_final['final_predicted'] = y_train_pred_final.Churn_Prob.map(
    ↪lambda x: 1 if x > 0.3 else 0)

y_train_pred_final.head()
```

```
[302]:
```

	Churn	Churn_Prob	CustID	predicted	0.0	0.1	0.2	0.3	0.4	0.5	0.6	\
0	0	0.204339	879	0	0	0	0	0	0	0	0	
1	0	0.215562	5790	0	0	0	0	0	0	0	0	
2	1	0.639615	6498	1	1	1	1	1	1	1	1	
3	1	0.687136	880	1	1	1	1	1	1	1	1	
4	1	0.735303	2784	1	1	1	1	1	1	1	1	

	0.7	0.8	0.9	Predicted_Label	final_predicted
0	0	0	0	0	0
1	0	0	0	0	0
2	1	1	1	0	1
3	1	1	1	0	1
4	1	1	1	0	1

```
[303]: # Let's check the overall accuracy.  
metrics.accuracy_score(y_train_pred_final.Churn, y_train_pred_final.  
    ↪final_predicted)
```

```
[303]: 0.76330759853718
```

```
[304]: confusion2 = metrics.confusion_matrix(y_train_pred_final.Churn, ↪  
    ↪y_train_pred_final.final_predicted )  
confusion2
```

```
[304]: array([[2751,  884],  
        [ 281, 1006]])
```

```
[305]: TP = confusion2[1,1] # true positive  
TN = confusion2[0,0] # true negatives  
FP = confusion2[0,1] # false positives  
FN = confusion2[1,0] # false negatives
```

```
[306]: # Let's see the sensitivity of our logistic regression model  
TP / float(TP+FN)
```

```
[306]: 0.7816627816627817
```

```
[307]: # Let us calculate specificity  
TN / float(TN+FP)
```

```
[307]: 0.756808803301238
```

```
[308]: # Calculate false postive rate - predicting churn when customer does not have ↪  
    ↪churned  
print(FP/ float(TN+FP))
```

```
0.24319119669876205
```

```
[309]: # Positive predictive value  
print (TP / float(TP+FP))
```

```
0.5322751322751322
```

```
[310]: # Negative predictive value  
print (TN / float(TN+ FN))
```

```
0.9073218997361477
```

```
[311]: # Preciosin and Recall  
#Looking at the confusion matrix again
```

```
confusion = metrics.confusion_matrix(y_train_pred_final.Churn,
    ↪y_train_pred_final.predicted )
confusion
```

```
[311]: array([[3277,  358],
           [ 614,  673]])
```

```
[281]: print("Precision =", confusion[1,1]/(confusion[0,1]+confusion[1,1]))
```

```
Precision = 0.6527643064985451
```

```
[282]: print("Recall =", confusion[1,1]/(confusion[1,0]+confusion[1,1]))
```

```
Recall = 0.5229215229215229
```

```
[283]: from sklearn.metrics import precision_score, recall_score
```

```
[284]: print("Precision =", precision_score(y_train_pred_final.Churn,
    ↪y_train_pred_final.predicted))
print("Recall =", recall_score(y_train_pred_final.Churn, y_train_pred_final.
    ↪predicted))
```

```
Precision = 0.6527643064985451
```

```
Recall = 0.5229215229215229
```

2.1.3 Precision and recall tradeoff

```
[312]: from sklearn.metrics import precision_recall_curve
y_train_pred_final.Churn, y_train_pred_final.predicted
```

```
[312]: (0      0
        1      0
        2      1
        3      1
        4      1
        ..
        4917    0
        4918    0
        4919    0
        4920    0
        4921    0
        Name: Churn, Length: 4922, dtype: int64,
        0      0
        1      0
        2      1
        3      1
        4      1)
```

```

..
4917    0
4918    0
4919    0
4920    0
4921    0
Name: predicted, Length: 4922, dtype: int64)

```

```


[313]: p , r , thresholds = precision_recall_curve(y_train_pred_final.Churn,
↪y_train_pred_final.predicted)
plt.plot(thresholds, p[:-1], "g-")
plt.plot(thresholds, r[:-1], "r-")

```

```

[313]: [<matplotlib.lines.Line2D at 0x794f512d85e0>]

```



2.1.4 Step 11: Making predictions on the test set

```

[286]: # test the data with y_test and X_test
X_test[['tenure', 'MonthlyCharges', 'TotalCharges']] = scaler.
↪transform(X_test[['tenure', 'MonthlyCharges', 'TotalCharges']])

```

```
[287]: X_test = X_test[col]
X_test.head()
```

```
[287]:      tenure  SeniorCitizen  Contract_One year  Contract_Two year \
942   -0.347623             0             False             False
3730   0.999203             0             False             False
1761   1.040015             0             False             True
2283  -1.286319             0             False             False
1872   0.346196             0             False             True

      PaymentMethod_Credit card (automatic)  PaymentMethod_Mailed check \
942                                     True             False
3730                                    True             False
1761                                    True             False
2283                                   False             True
1872                                   False             False

      InternetService_Fiber optic  InternetService_No  MultipleLines_Yes \
942                             True             False             False
3730                             True             False             True
1761                             False            True             True
2283                             True             False             False
1872                             False            True             False

      StreamingTV_Yes  StreamingMovies_Yes
942                False                True
3730                True                True
1761                False                False
2283                False                False
1872                False                False
```

```
[288]: y_test.head()
```

```
[288]: 942      0
3730    1
1761    0
2283    1
1872    0
Name: Churn, dtype: int64
```

```
[289]: X_test_sm.head()
```

```
[289]:      const  tenure  SeniorCitizen  Contract_One year  Contract_Two year \
942     1.0 -0.347623             0             0             0
3730     1.0  0.999203             0             0             0
1761     1.0  1.040015             0             0             1
2283     1.0 -1.286319             0             0             0
```

1872	1.0	0.346196	0	0	1
------	-----	----------	---	---	---

	PaymentMethod_Credit card (automatic)	PaymentMethod_Mailed check	\
942	1	0	
3730	1	0	
1761	1	0	
2283	0	1	
1872	0	0	

	InternetService_Fiber optic	InternetService_No	MultipleLines_Yes	\
942	1	0	0	
3730	1	0	1	
1761	0	1	1	
2283	1	0	0	
1872	0	1	0	

	StreamingTV_Yes	StreamingMovies_Yes
942	0	1
3730	1	1
1761	0	0
2283	0	0
1872	0	0

```
[290]: X_test_sm.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 2110 entries, 942 to 4987
Data columns (total 12 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   const                                     2110 non-null   float64
1   tenure                                   2110 non-null   float64
2   SeniorCitizen                             2110 non-null   int64
3   Contract_One year                         2110 non-null   int64
4   Contract_Two year                         2110 non-null   int64
5   PaymentMethod_Credit card (automatic)    2110 non-null   int64
6   PaymentMethod_Mailed check                2110 non-null   int64
7   InternetService_Fiber optic               2110 non-null   int64
8   InternetService_No                        2110 non-null   int64
9   MultipleLines_Yes                         2110 non-null   int64
10  StreamingTV_Yes                           2110 non-null   int64
11  StreamingMovies_Yes                       2110 non-null   int64
dtypes: float64(2), int64(10)
memory usage: 214.3 KB
```

```
[291]: bool_columns = X_test.select_dtypes(include=['bool']).columns
X_test[bool_columns] = X_test[bool_columns].astype(int)
```

```
[292]: X_test_sm = sm.add_constant(X_test[col])
```

```
[293]: logm_test = sm.GLM(y_test, X_test_sm, family = sm.families.Binomial())
res = logm_test.fit()
print(res.summary())
```

Generalized Linear Model Regression Results

```
=====
Dep. Variable:          Churn    No. Observations:          2110
Model:                  GLM      Df Residuals:              2098
Model Family:           Binomial  Df Model:                  11
Link Function:          Logit     Scale:                    1.0000
Method:                 IRLS      Log-Likelihood:           -925.02
Date:                   Wed, 07 Aug 2024    Deviance:                 1850.0
Time:                   03:42:27    Pearson chi2:             1.98e+03
No. Iterations:         7          Pseudo R-squ. (CS):       0.2600
Covariance Type:        nonrobust
=====
```

```
=====
                                coef    std err          z
P>|z|      [0.025    0.975]
-----
const                                -1.5693    0.162    -9.709
0.000    -1.886    -1.252
tenure                                -0.8892    0.095    -9.410
0.000    -1.074    -0.704
SeniorCitizen                        -0.0548    0.149    -0.368
0.713    -0.346     0.237
Contract_One year                     -0.8690    0.190    -4.571
0.000    -1.242    -0.496
Contract_Two year                     -1.9106    0.310    -6.162
0.000    -2.518    -1.303
PaymentMethod_Credit card (automatic) -0.2058    0.164    -1.252
0.211    -0.528     0.116
PaymentMethod_Mailed check            -0.1075    0.163    -0.661
0.509    -0.427     0.212
InternetService_Fiber optic           0.9304    0.141     6.580
0.000     0.653     1.208
InternetService_No                    -0.5142    0.213    -2.412
0.016    -0.932    -0.096
MultipleLines_Yes                     0.3264    0.136     2.405
0.016     0.060     0.592
StreamingTV_Yes                       0.1900    0.141     1.347
0.178    -0.087     0.467
StreamingMovies_Yes                   0.5022    0.142     3.526
0.000     0.223     0.781
=====
```

```
=====
=====
```

```
[294]: y_test_pred = res.predict(X_test_sm)
```

```
[295]: y_test_pred[:10]
```

```
[295]: 942      0.491669
      3730    0.328609
      1761    0.008177
      2283    0.598074
      1872    0.013362
      1970    0.617101
      2532    0.235379
      1616    0.007807
      2485    0.588950
      5914    0.214762
      dtype: float64
```

```
[314]: # Converting y_pred to a dataframe which is an array
      y_pred_1 = pd.DataFrame(y_test_pred)
      y_pred_1
```

```
[314]:      0
      942  0.491669
      3730 0.328609
      1761 0.008177
      2283 0.598074
      1872 0.013362
      ...
      1289 0.038491
      3508 0.052596
      6765 0.006749
      3598 0.382200
      4987 0.006279
```

```
[2110 rows x 1 columns]
```

```
[316]: # Converting y_test to dataframe
      y_test_df = pd.DataFrame(y_test)
```

```
[317]: # Putting CustID to index
      y_test_df['CustID'] = y_test_df.index
```

```
[318]: # Removing index for both dataframes to append them side by side
      y_pred_1.reset_index(drop=True, inplace=True)
      y_test_df.reset_index(drop=True, inplace=True)
```



```
[319]: # Appending y_test_df and y_pred_1
y_pred_final = pd.concat([y_test_df, y_pred_1],axis=1)
```

```
[320]: y_pred_final.head()
```

```
[320]:
```

	Churn	CustID	0
0	0	942	0.491669
1	1	3730	0.328609
2	0	1761	0.008177
3	1	2283	0.598074
4	0	1872	0.013362

```
[321]: # Renaming the column
y_pred_final= y_pred_final.rename(columns={ 0 : 'Churn_Prob'})
```

```
[324]: # Rearranging the columns

y_pred_final = y_pred_final.reindex(columns=['CustID', 'Churn', 'Churn_Prob'])

# Let's see the head of y_pred_final
y_pred_final.head()
```

```
[324]:
```

	CustID	Churn	Churn_Prob
0	942	0	0.491669
1	3730	1	0.328609
2	1761	0	0.008177
3	2283	1	0.598074
4	1872	0	0.013362

```
[326]: # Base don the precision and Recall trade off graph choosing cutoff as 0.82 ,
↳ please use accordingly for your data
y_pred_final['final_predicted'] = y_pred_final.Churn_Prob.map(lambda x: 1 if x
↳ > 0.82 else 0)
```

The choice of a cutoff value, such as 0.82, is often based on the specific context and goals of the classification problem. In binary classification, the probability cutoff determines the threshold above which a prediction is considered positive (e.g., churn) and below which it is considered negative (e.g., no churn). The default threshold is typically 0.5, but there are several reasons why you might choose a different threshold like 0.82:

Class Imbalance: If the classes are imbalanced (e.g., there are many more non-churners than churners), a lower threshold might help in identifying more positive cases, improving sensitivity/recall at the cost of specificity.

Cost-Benefit Analysis: The costs of false positives and false negatives might differ significantly. For example, if false negatives (not identifying a churner) are more costly than false positives (incorrectly identifying a non-churner as a churner), you might lower the threshold to reduce false negatives.

Optimization for Specific Metrics: Depending on the business requirements, you might optimize for metrics such as F1 score, sensitivity (recall), or precision. A threshold like 0.82 might have been found to optimize these metrics during model validation.

Receiver Operating Characteristic (ROC) Curve: The ROC curve plots the true positive rate (sensitivity) against the false positive rate (1-specificity) for various threshold values. The point closest to the top-left corner of the ROC space (where both sensitivity and specificity are high) can be chosen as the optimal threshold. This process might lead to a threshold like 0.82.

Precision-Recall Tradeoff: In some cases, particularly with imbalanced datasets, a Precision-Recall curve might be used to find the threshold that provides the best tradeoff between precision and recall.

```
[327]: confusion2 = metrics.confusion_matrix(y_pred_final.C churn, y_pred_final.  
      ↪ final_predicted )  
      confusion2
```

```
[327]: array([[1528,    0],  
      [ 580,    2]])
```

```
[328]: TP = confusion2[1,1] # true positive  
      TN = confusion2[0,0] # true negatives  
      FP = confusion2[0,1] # false positives  
      FN = confusion2[1,0] # false negatives
```

```
[329]: # Let's see the sensitivity of our logistic regression model  
      TP / float(TP+FN)
```

```
[329]: 0.003436426116838488
```

```
[330]: # Let us calculate specificity  
      TN / float(TN+FP)
```

```
[330]: 1.0
```

```
[ ]:
```